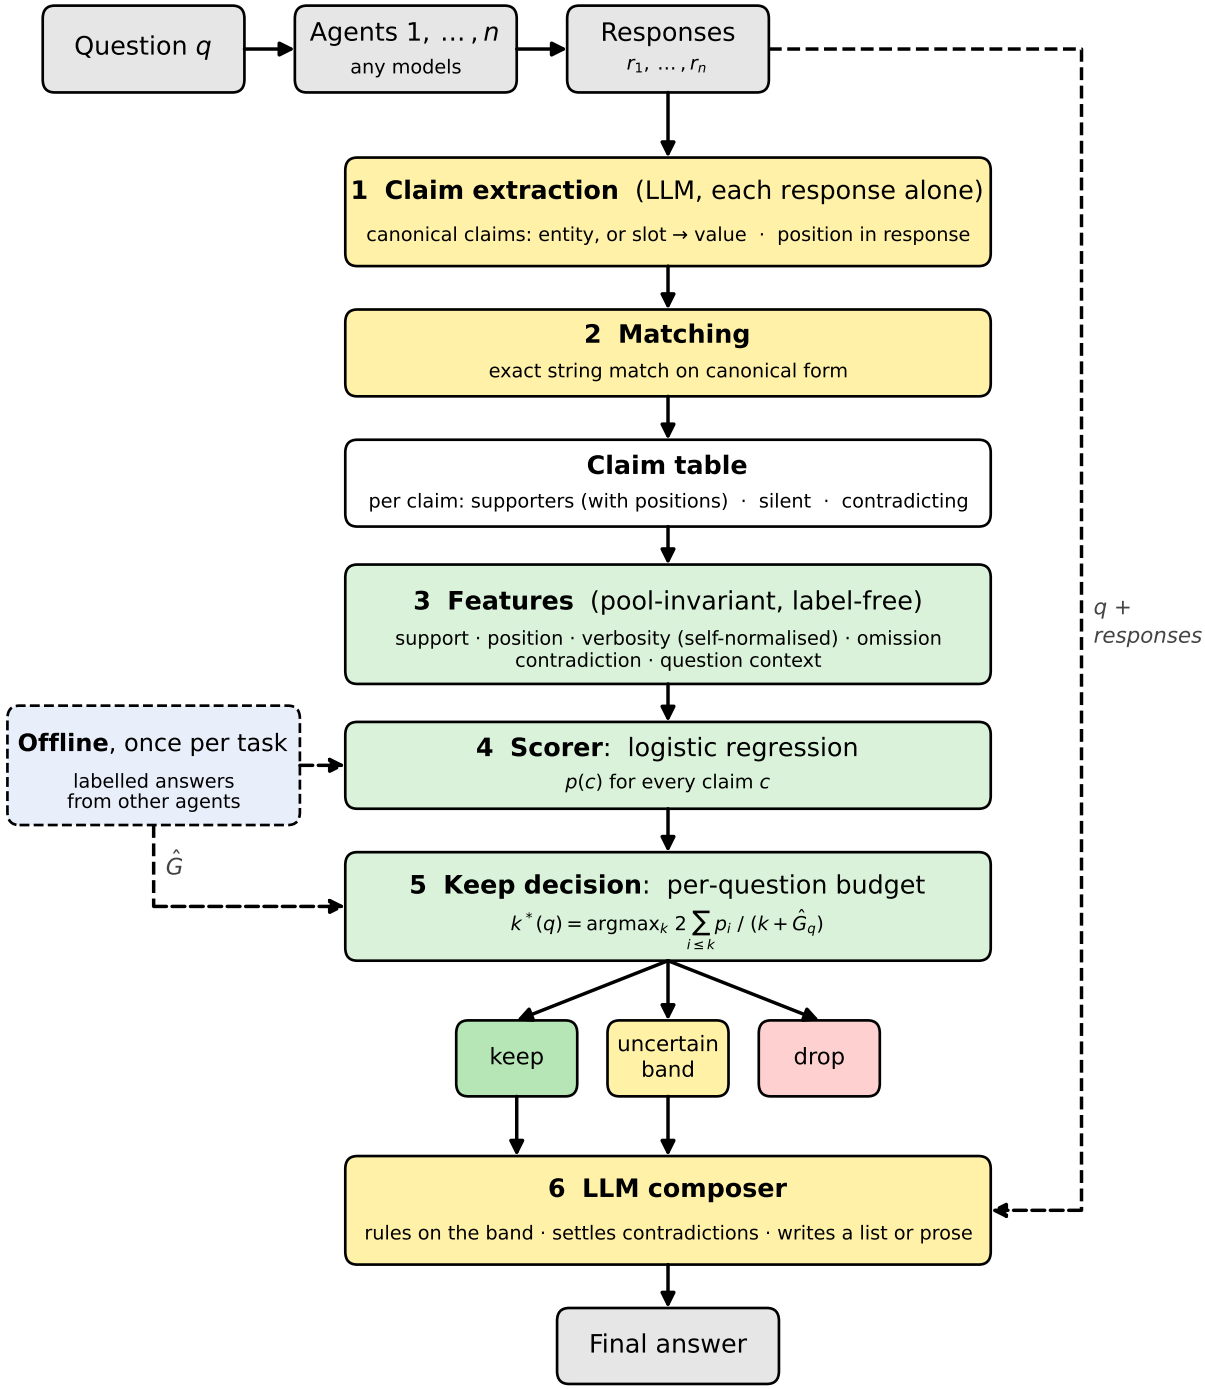
Question q
Agents 1, … , n
any models
Responses
r_1, … , r_n
1 Claim extraction (LLM, each response alone)
canonical claims: entity, or slot → value · position in response
2 Matching
exact string match on canonical form
Claim table
per claim: supporters (with positions) · silent · contradicting
3 Features (pool-invariant, label-free)
support · position · verbosity (self-normalised) · omission
contradiction · question context
q + responses
Offline, once per task
labelled answers from other agents
4 Scorer: logistic regression
$p(c)$ for every claim c
$\hat{G}$
5 Keep decision: per-question budget
$k^*(q) = \text{argmax}_k \, 2 \sum_{i \le k} p_i \,/\, (k + \hat{G}_q)$
keep
uncertain band
drop
6 LLM composer
rules on the band · settles contradictions · writes a list or prose
Final answer
run on real models
built, not yet run
offline